

CSC 112: Computer Operating Systems

Lecture 2

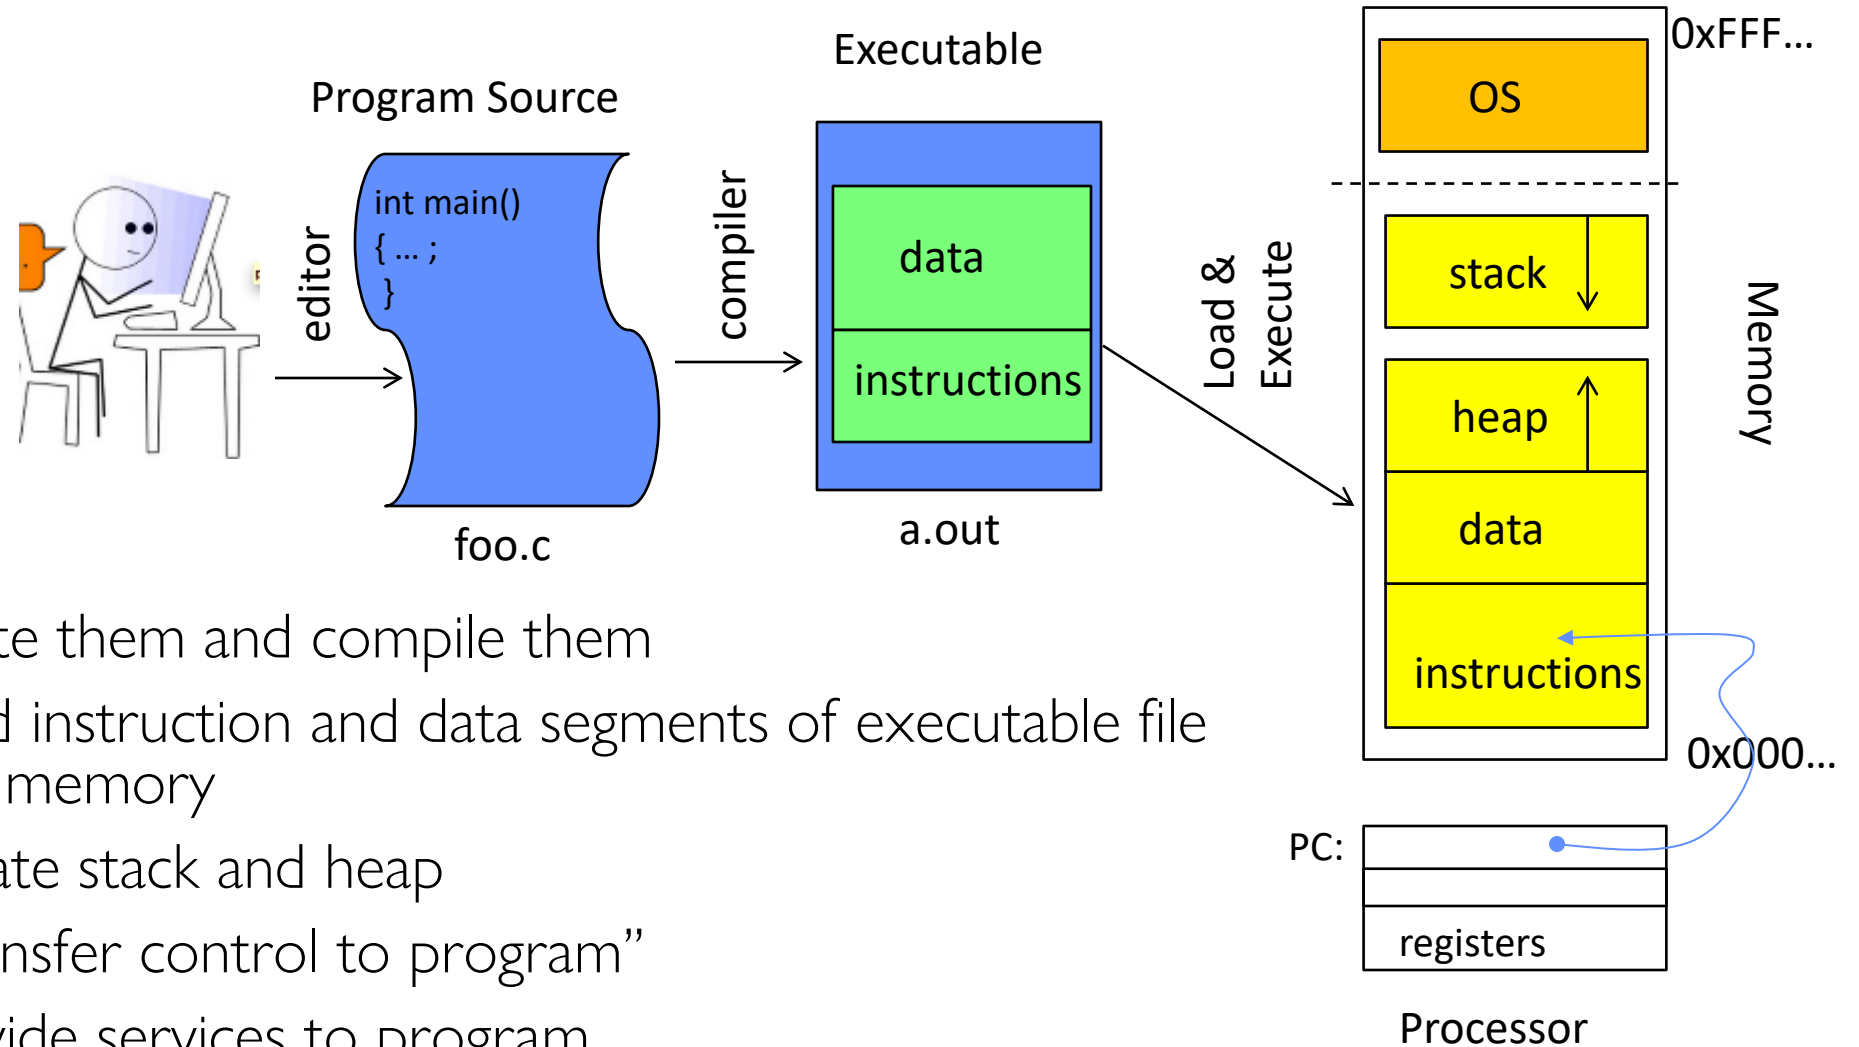
Four Fundamental OS Concepts

Department of Computer Science,
Hofstra University

Today: Four Fundamental OS Concepts

- **Thread: Execution Context**
 - Fully describes program state
 - Program Counter, Registers, Execution Flags, Stack
- **Address space (with or w/o translation)**
 - Set of memory addresses accessible to program (for read or write)
 - May be distinct from memory space of the physical machine (in which case programs operate in a virtual address space)
- **Process: an instance of a running program**
 - Protected Address Space + One or more Threads
- **Dual mode operation / Protection**
 - Only the “system” has the ability to access certain resources
 - Combined with translation, isolates programs from each other and the OS from programs

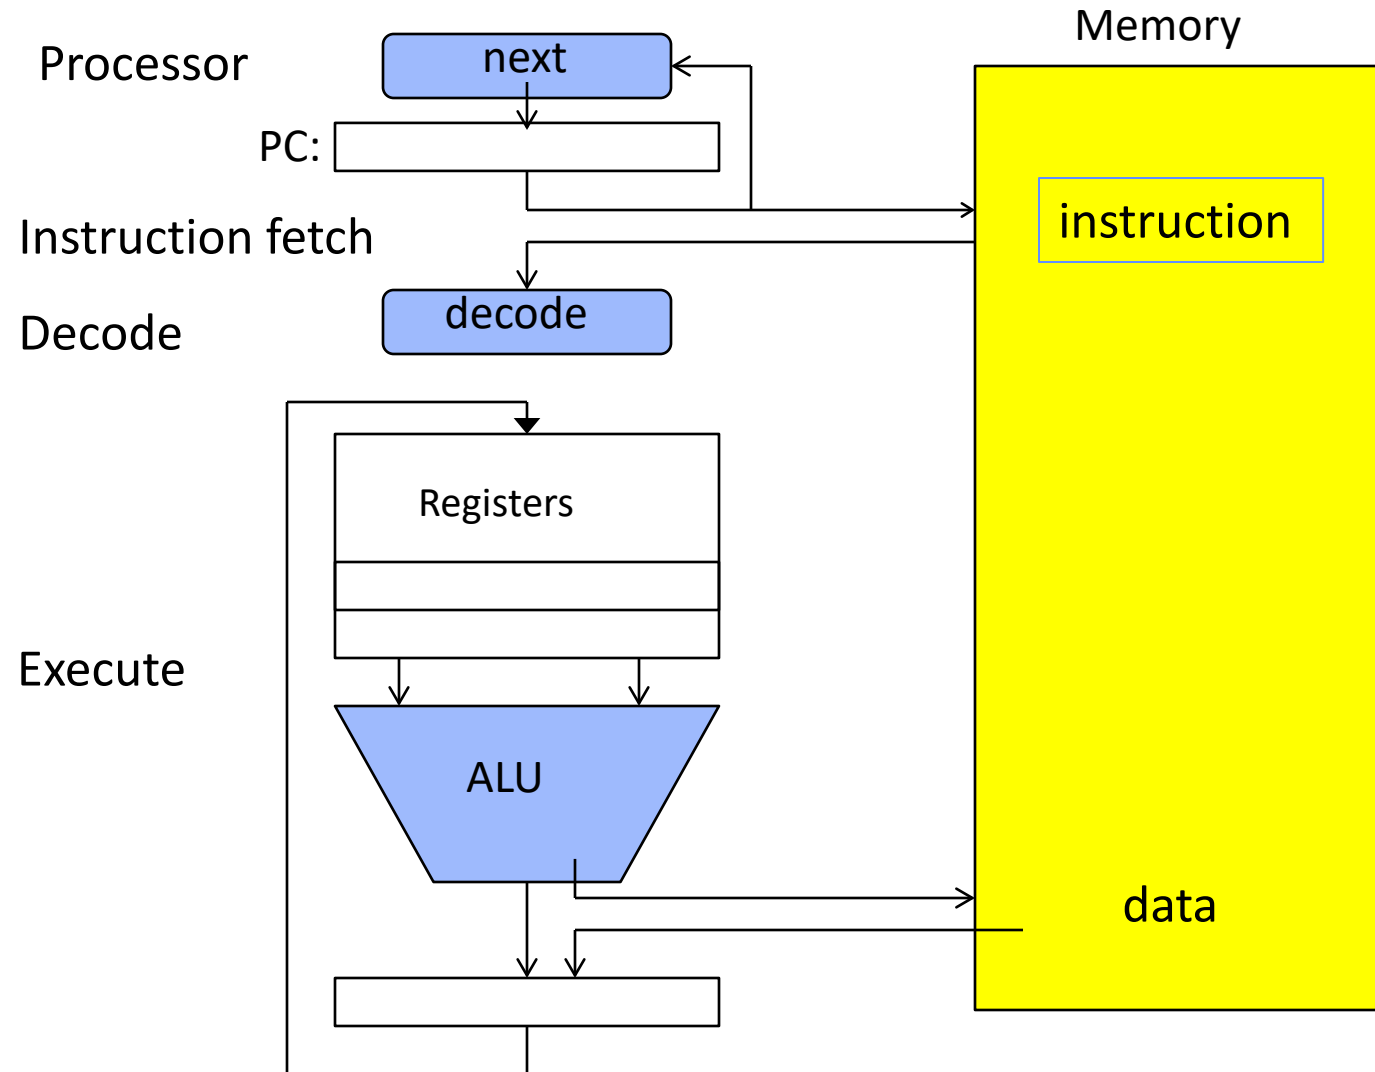
OS Bottom Line: Run Programs



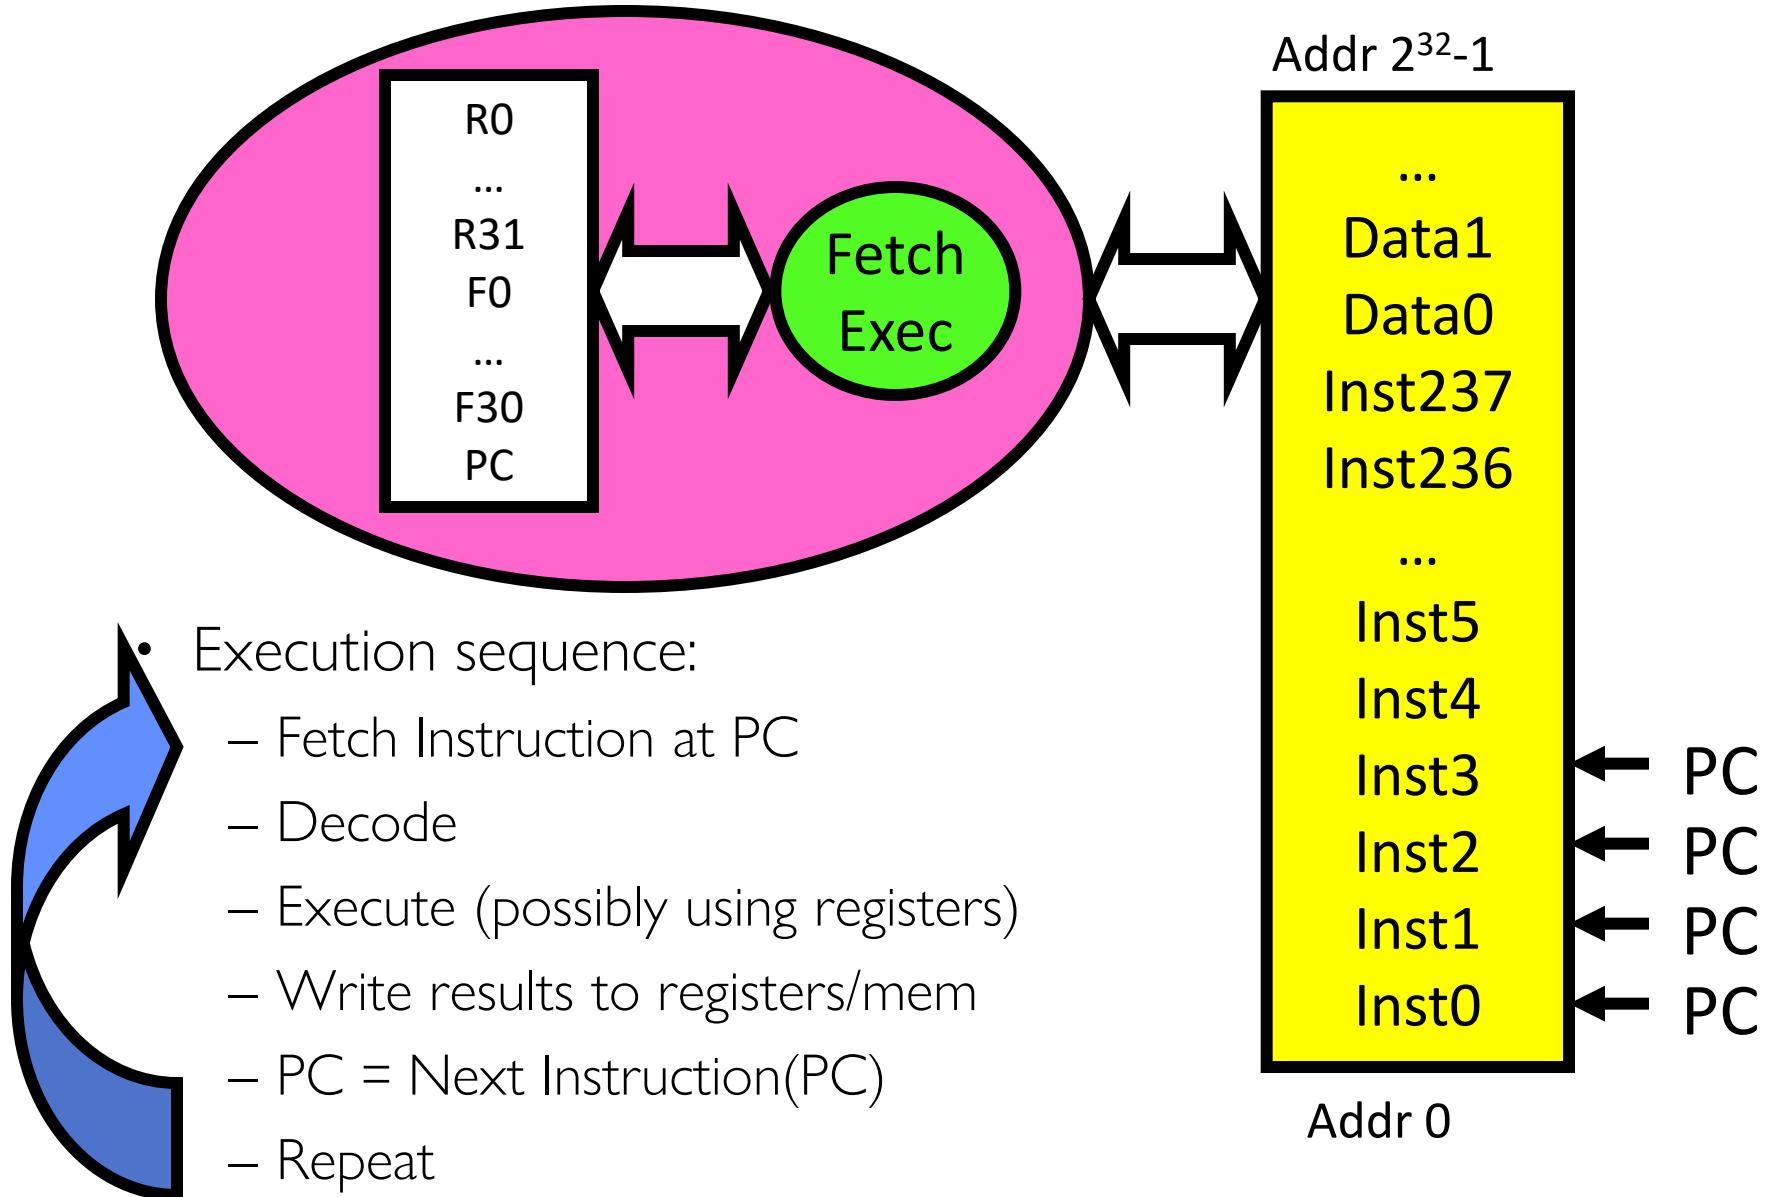
- Write them and compile them
- Load instruction and data segments of executable file into memory
- Create stack and heap
- “Transfer control to program”
- Provide services to program
- While protecting OS and program

Instruction Fetch/Decode/Execute

The instruction cycle




What happens during program execution?




First OS Concept: Thread of Control

- **Thread:** Single unique execution context
 - Program Counter, Registers, Execution Flags, Stack, Memory State
- A thread is *executing* on a processor (core) when it is *resident* in the processor registers
- Resident means: Registers hold the root state (context) of the thread:
 - Including program counter (PC) register & currently executing instruction
 - » PC points at next instruction *in memory*
 - » Instructions stored in memory
 - Including intermediate values for ongoing computations
 - » Can include actual values (like integers) or pointers to values *in memory*
 - Stack pointer holds the address of the top of stack (which is *in memory*)
 - The rest is “*in memory*”
- A thread is *suspended* (not *executing*) when its state *is not* loaded (resident) into the processor

Illusion of Multiple Processors




Programmer's View

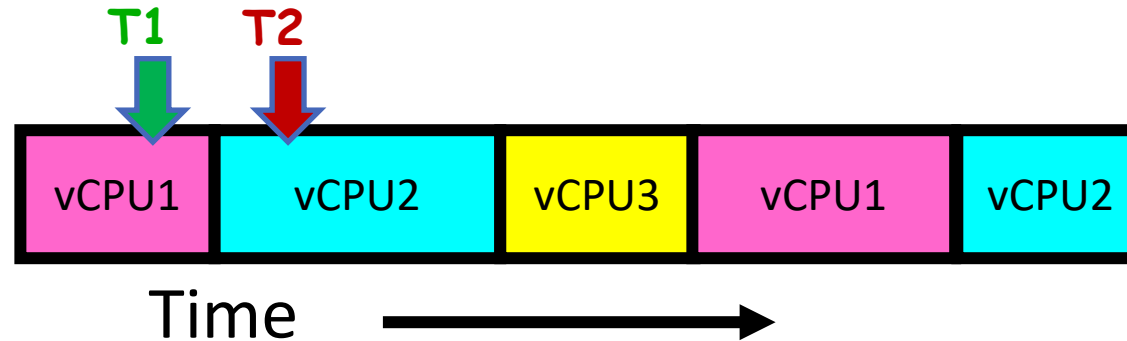


- Assume a single processor (core). How do we provide the illusion of multiple processors?
 - Multiplex in time!
- Threads are *virtual cores*
- Contents of virtual core (thread):
 - Program counter, stack pointer
 - Registers
- Where is “it” (the thread)?
 - On the real (physical) core, or
 - Saved in chunk of memory – called the *Thread Control Block (TCB)*

Illusion of Multiple Processors (Continued)




Programmer's View

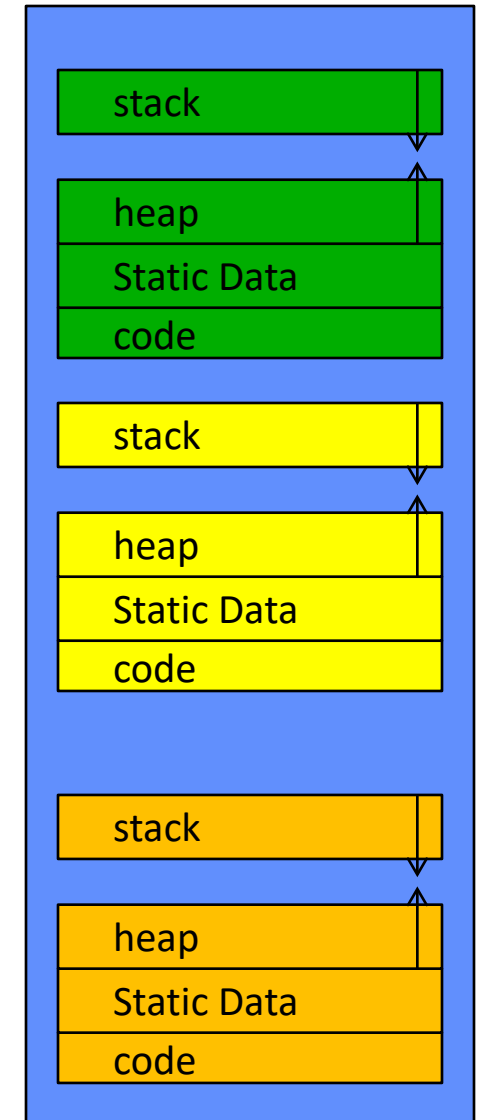


- Consider:
 - At T1: vCPU1 on real core, vCPU2 in memory
 - At T2: vCPU2 on real core, vCPU1 in memory
- What happened?
 - OS ran
 - Saved program counter (PC), stack pointer (SP), ... in vCPU1's TCB (in memory)
 - Loaded PC, SP, ... from vCPU2's TCB, jumped to PC
- What triggered this switch?
 - Timer, voluntary yield, I/O...

Multiprogramming - Multiple Threads of Control

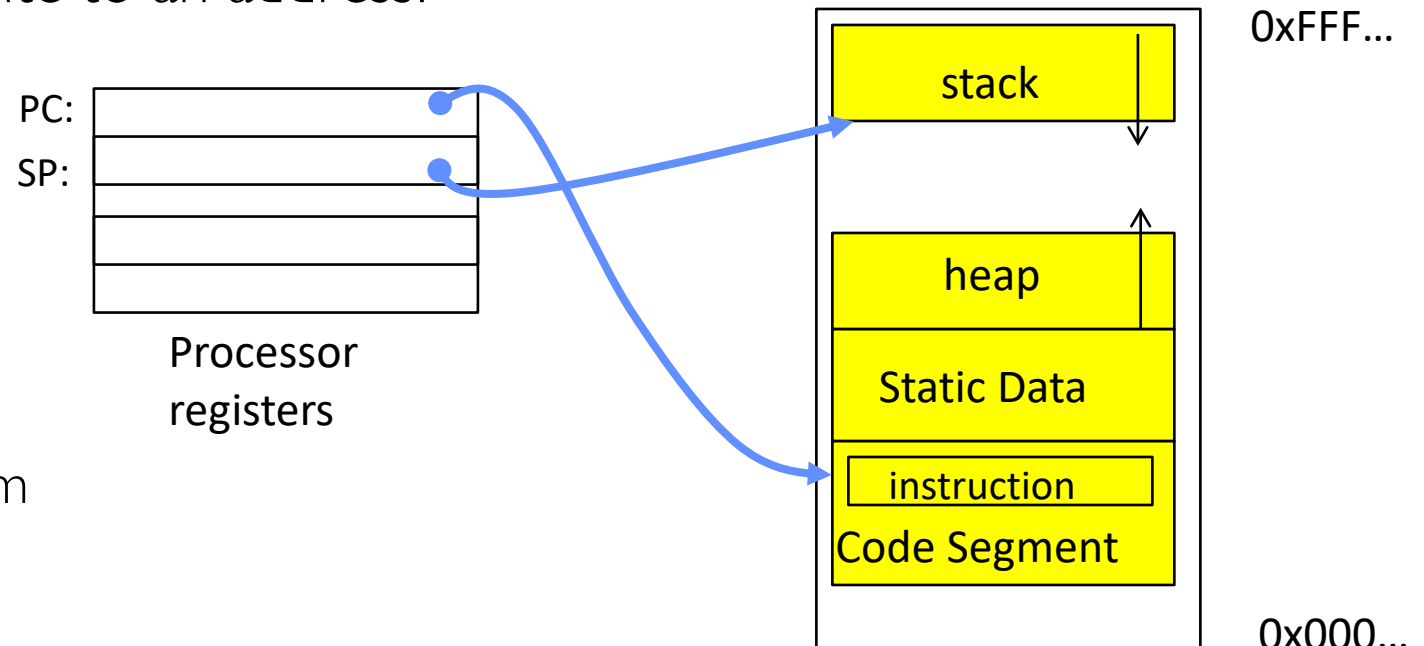


- Thread Control Block (TCB)
 - Holds contents of registers when thread not running
- Where are TCBs stored?
 - In the kernel




Second OS Concept: Address Space

- Address space \Rightarrow the set of accessible addresses + state associated with them:
 - For 32-bit processor: $2^{32} = 4$ billion ($\sim 10^9$) addresses
 - For 64-bit processor: $2^{64} = 18$ quintillion ($\sim 10^{18}$) addresses
- What happens when you read or write to an address?
 - Perhaps acts like regular memory
 - Perhaps ignores writes
 - Perhaps causes I/O operation
 - » (Memory-mapped I/O)
 - Perhaps causes exception (fault)
 - Communicates with another program
 -



Previous discussion of threads: Very Simple Multiprogramming

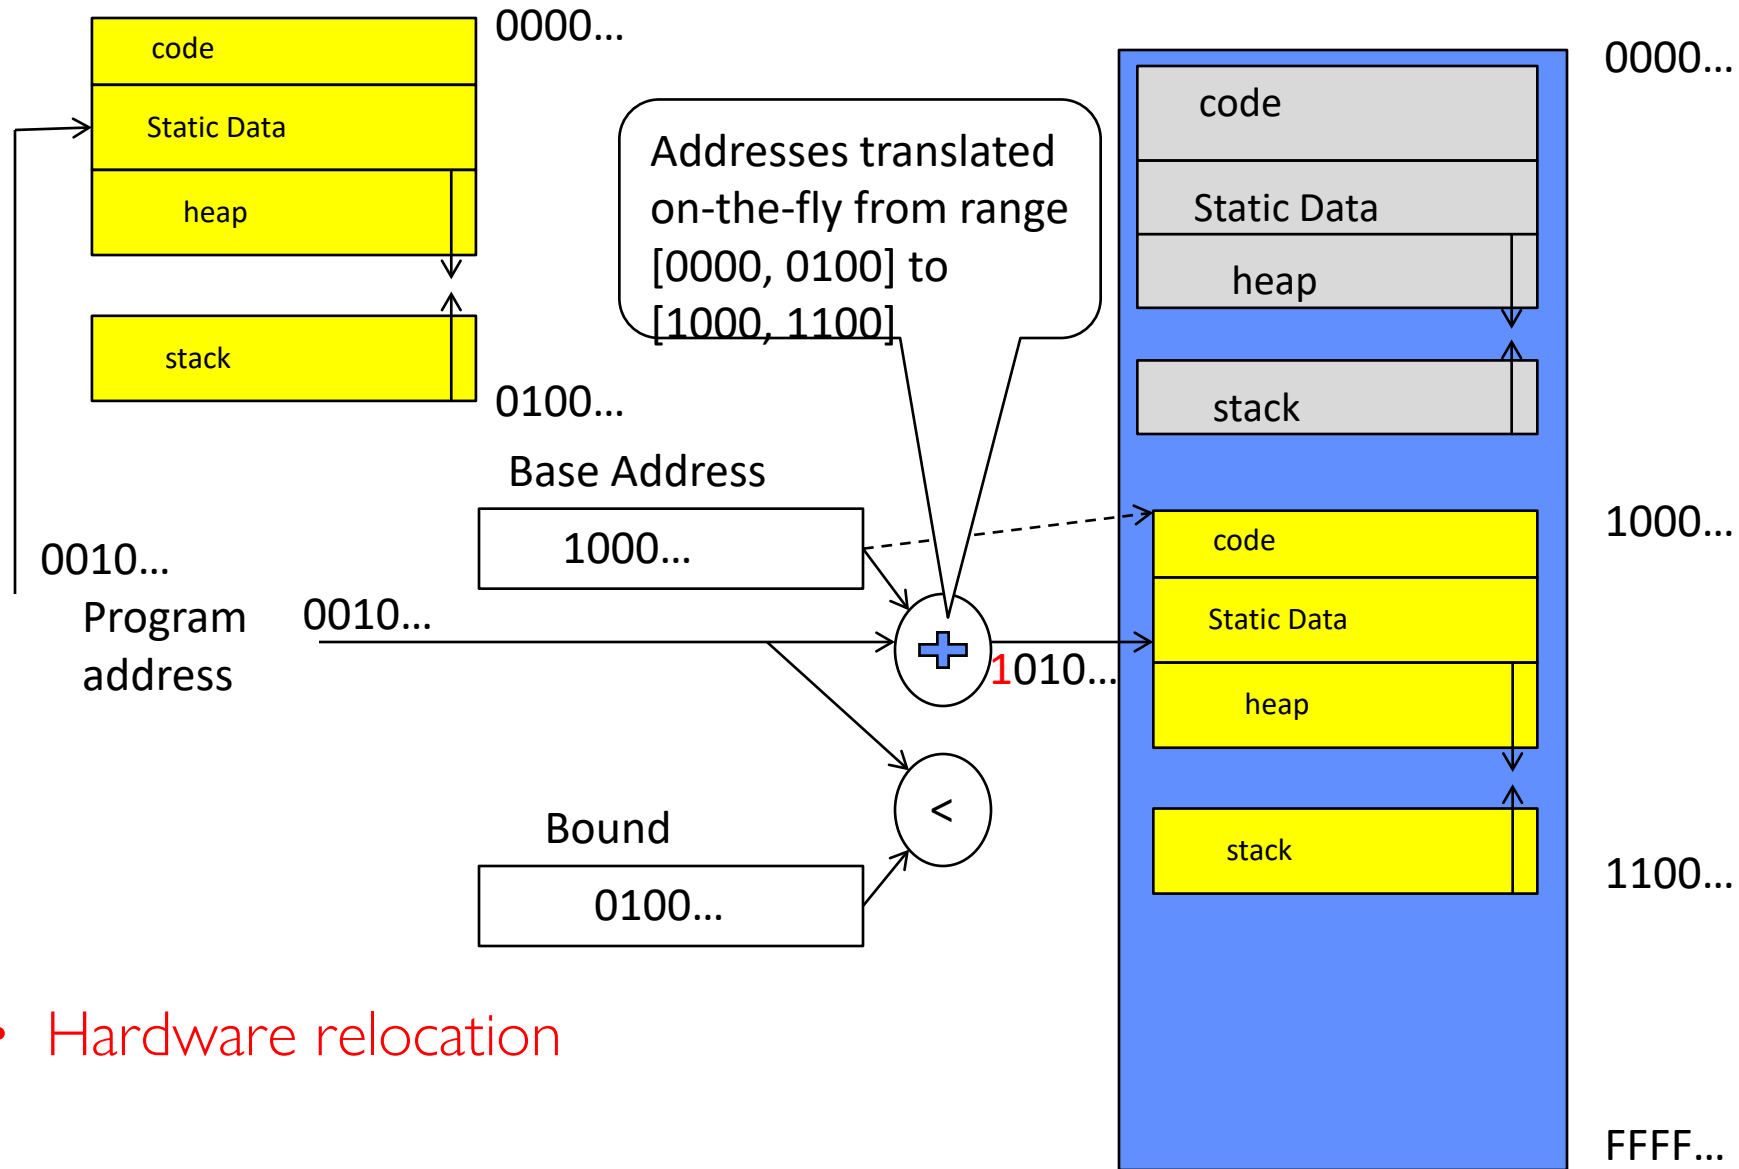
- All vCPU's share non-CPU resources
 - Memory, I/O Devices
- Each thread can read/write memory
 - Perhaps data of others
 - can overwrite OS ?
- This approach is used in
 - Very early days of computing
 - Embedded applications
 - MacOS 1-9/Windows 3.1 (switch only with voluntary yield)
 - Windows 95-ME (switch with yield or timer)
- However it is risky...



Simple Multiplexing has no Protection!

- Operating System must protect itself from user programs
 - Reliability: compromising the operating system generally causes it to crash
 - Security: limit the scope of what threads can do
 - Privacy: limit each thread to the data it is permitted to access
 - Fairness: each thread should be limited to its appropriate share of system resources (CPU time, memory, I/O, etc)
- OS must protect User programs from one another
 - Prevent threads owned by one user from impacting threads owned by another user
 - Example: prevent one user from stealing secret information from another user

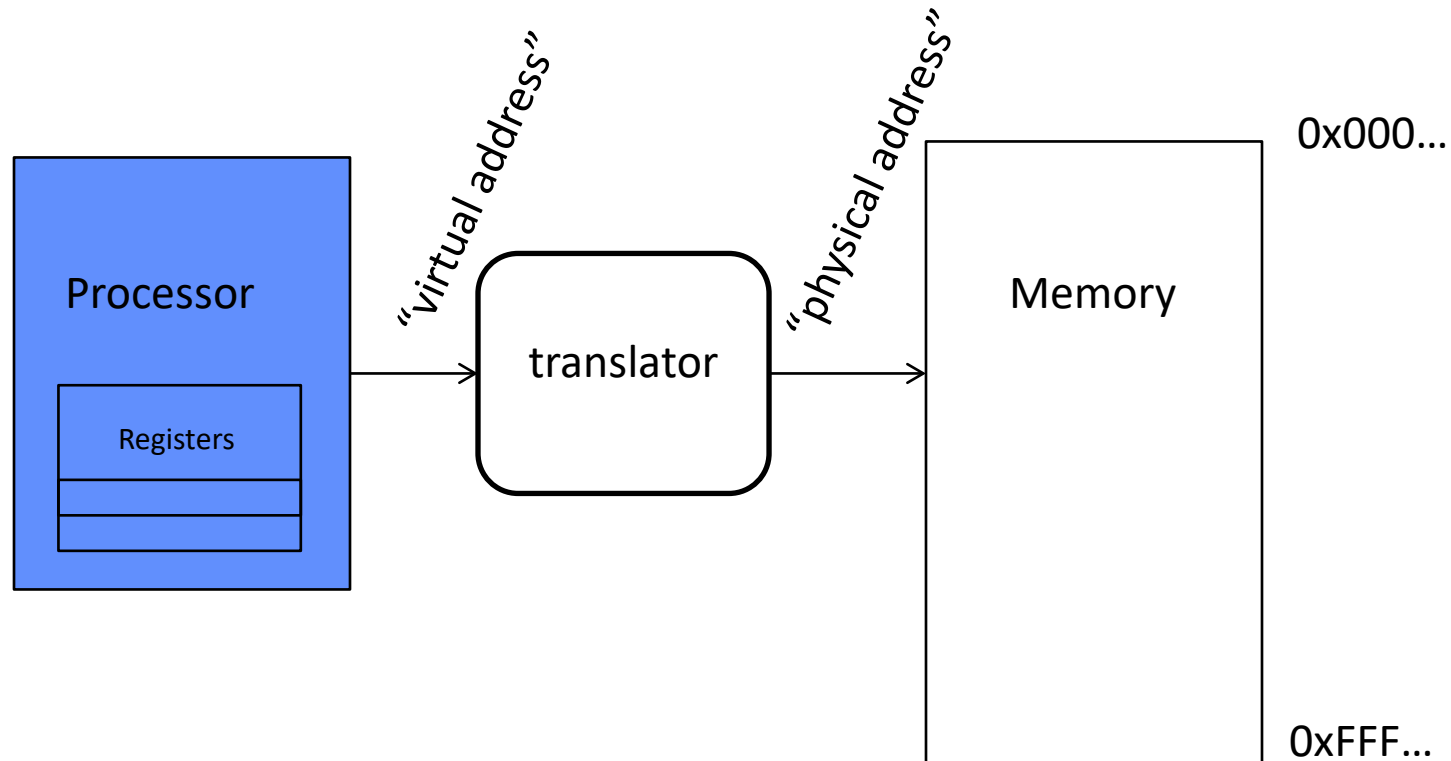
Simple address translation with Base and Bound



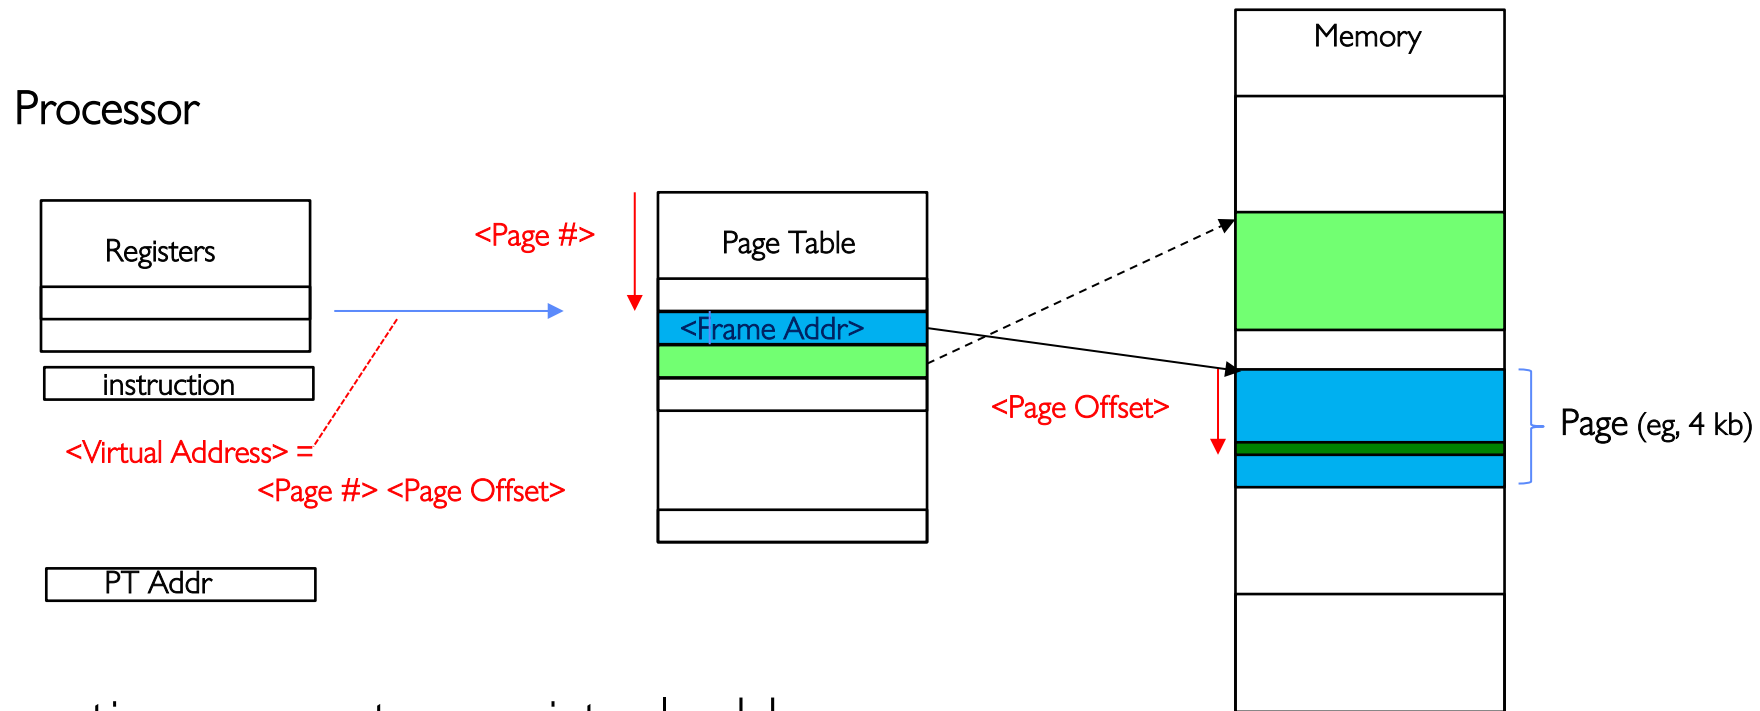
- Hardware relocation

Virtual-to-Physical Address Space Translation

- Program operates in a virtual address space that is distinct from the physical memory space of the machine
 - Break the entire virtual address space into equal size chunks (i.e., pages)
- Hardware translates address using a **page table**
 - Special hardware register stores pointer to page table
 - Treat memory as page size frames and put any page into any frame ...



Paged Virtual Address

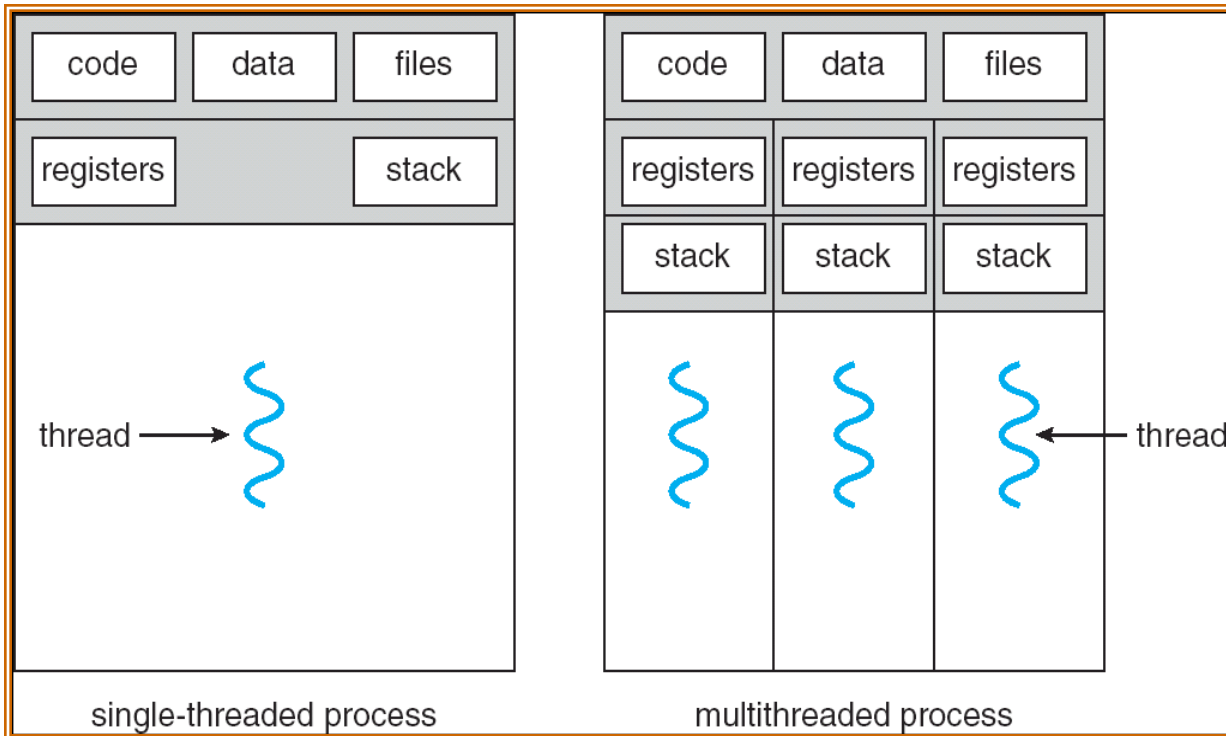


- Instructions operate on virtual addresses
 - Instruction address, load/store data address
- Translated to a physical address through a Page Table by the hardware
- Any Page of address space can be in any (page sized) frame in memory
 - Or not-present (access generates a page fault)

Third OS Concept: Process

- **Definition:** execution environment with Restricted Rights
 - **(Protected) Address Space with One or More Threads**
 - Owns memory (address space)
 - Owns file descriptors, file system context, ...
 - Encapsulate one or more threads sharing process resources
- Application program executes as a process
 - Complex applications can fork/exec child processes
- Why **processes**?
 - Protected from each other!
 - OS Protected from them
 - Processes provides memory protection
- Fundamental tradeoff between protection and efficiency
 - Communication easier *within* a process
 - Communication harder *between* processes

Single and Multithreaded Processes



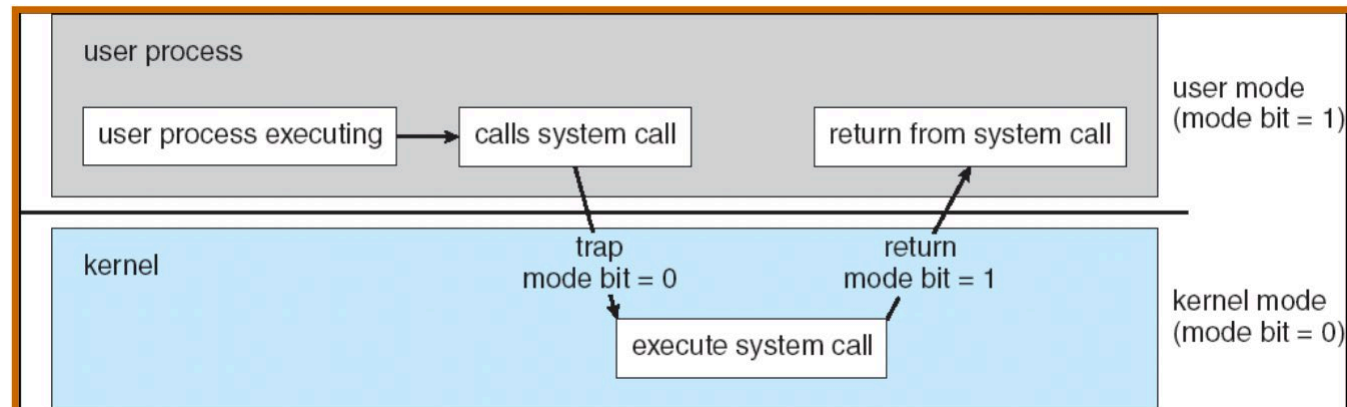
- Threads encapsulate **concurrency**:
 - “Active” component
- Address spaces encapsulate **protection**:
 - “Passive” component
 - Keeps buggy programs from crashing the system
- Why have multiple threads per address space?
 - Parallelism: take advantage of actual hardware parallelism (e.g. multicore)
 - Concurrency: ease of handling I/O and other simultaneous events

Protection and Isolation


- Why Do We Need Processes??
 - Reliability: bugs can only overwrite memory of process they are in
 - Security and privacy: malicious or compromised process can't read or write other process' data
 - Fairness: enforce shares of disk, CPU
- Mechanisms:
 - Address translation: address space only contains its own data
 - BUT: why can't a process change the page table pointer?
 - » Or use I/O instructions to bypass the system?
 - Hardware must support **privilege levels**

Fourth OS Concept: Dual Mode Operation


- **Hardware** provides at least two modes (at least 1 mode bit):
 1. **Kernel Mode** (or “supervisor” mode)
 2. **User Mode**
- Certain operations are **prohibited** when running in user mode
 - Changing the page table pointer, disabling interrupts, interacting directly w/ hardware, writing to kernel memory
- Carefully controlled transitions between user mode and kernel mode
 - System calls, interrupts, exceptions



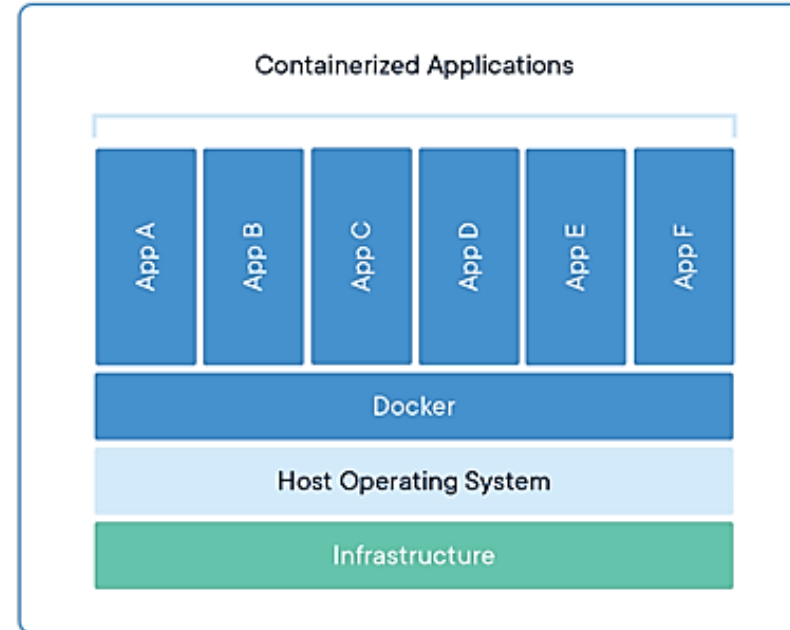
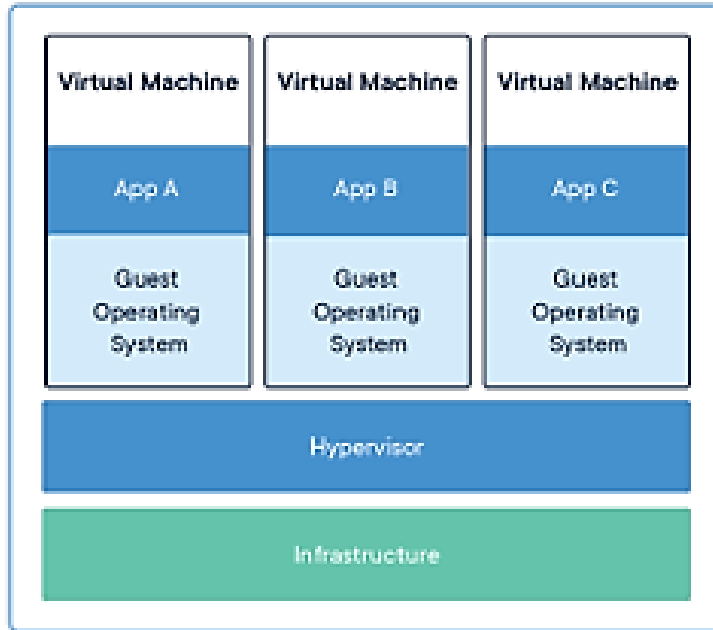
For example: UNIX System Structure



User/Kernel (Privileged) Mode



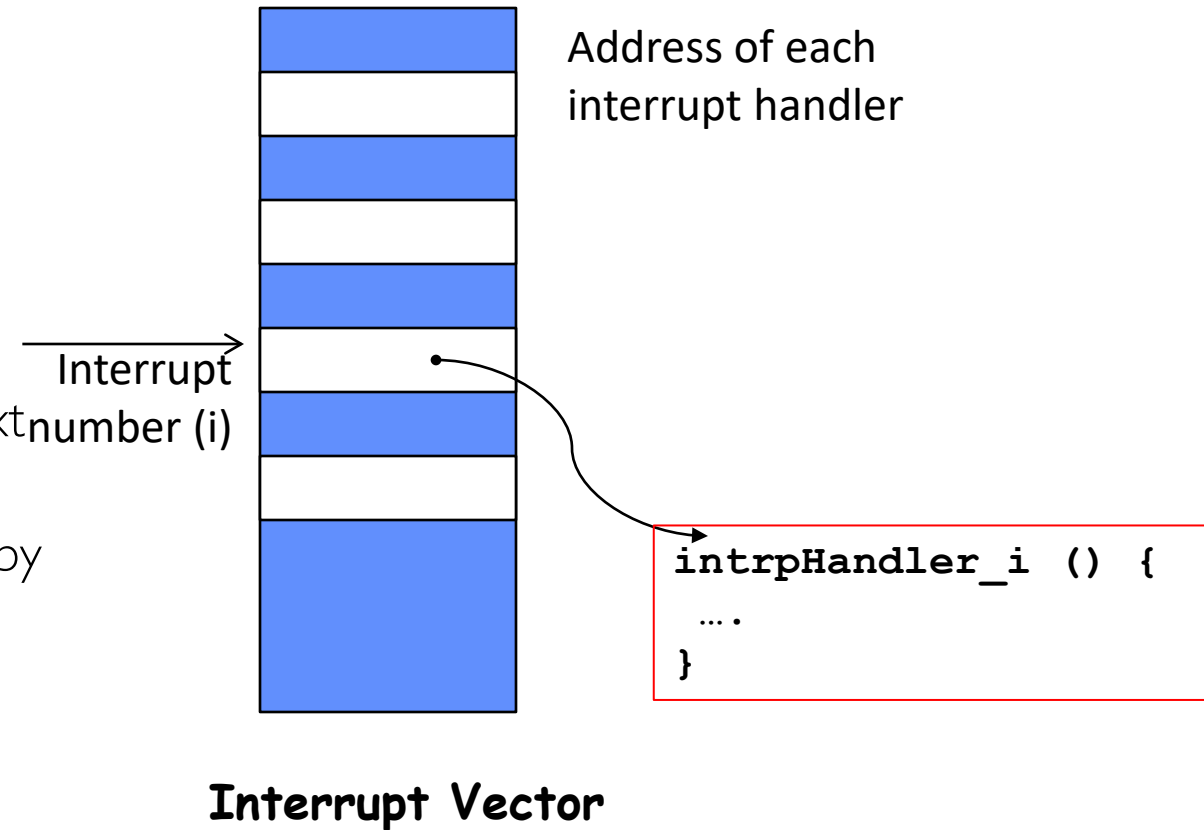
Additional Layers of Protection for Modern Systems



- Additional layers of protection through virtual machines or containers
 - Run a complete operating system in a virtual machine
 - Package all the libraries associated with an app into a container for execution
- More on these ideas later in the class

3 types of User \Rightarrow Kernel Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Put the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - » e. g., Timer, I/O device
 - Independent of user process
- Trap or Exception
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go?

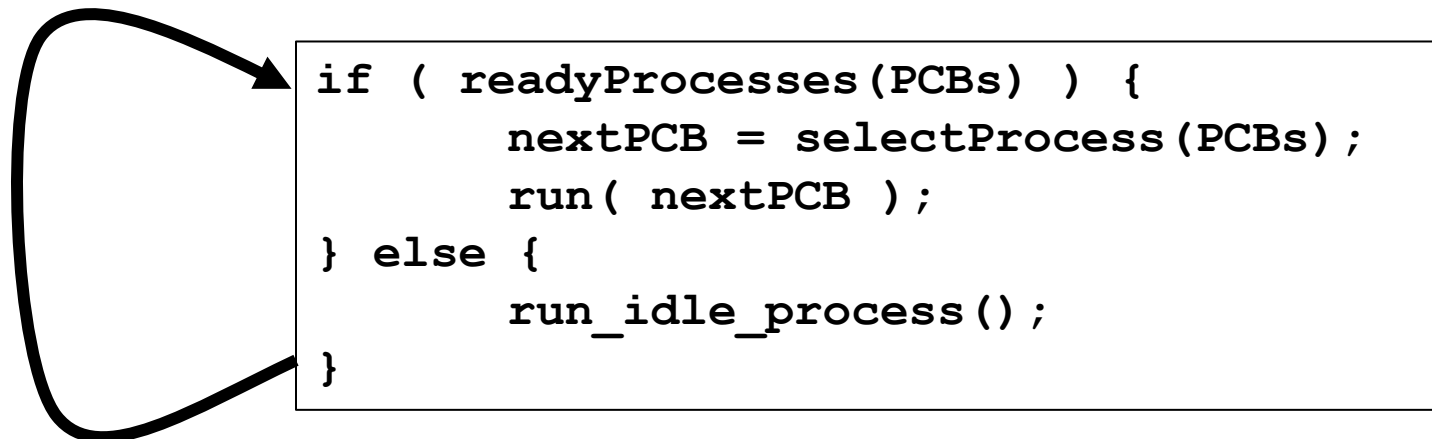


Running Many Programs ???

- We have the basic mechanism to
 - switch between user processes and the kernel,
 - the kernel can switch among user processes,
 - Protect OS from user processes and processes from each other
- Questions ???
- How do we decide which user process to run?
- How do we represent user processes in the OS?
- How do we pack up the process and set it aside?
- How do we get a stack and heap for the kernel?
- ...

Process Control Block and Scheduler

- Kernel represents each process as a process control block (PCB)
 - Status (running, ready, blocked, ...)
 - Register state (when not ready)
 - Process ID (PID), User, Executable, Priority, ...
 - Execution time, ...
 - Memory space, translation, ...
- Kernel Scheduler maintains a data structure containing the PCBs
 - Scheduling algorithm selects the next one to run



Conclusion: Four Fundamental OS Concepts

- **Thread: Execution Context**
 - Fully describes program state
 - Program Counter, Registers, Execution Flags, Stack
- **Address space (with or w/o translation)**
 - Set of memory addresses accessible to program (for read or write)
 - Virtual address space distinct from physical address space of the machine
- **Process: an instance of a running program**
 - Protected Address Space + One or more Threads
- **Dual mode operation / Protection**
 - Only the “system” has the ability to access certain resources
 - Combined with translation, isolates programs from each other and the OS from programs